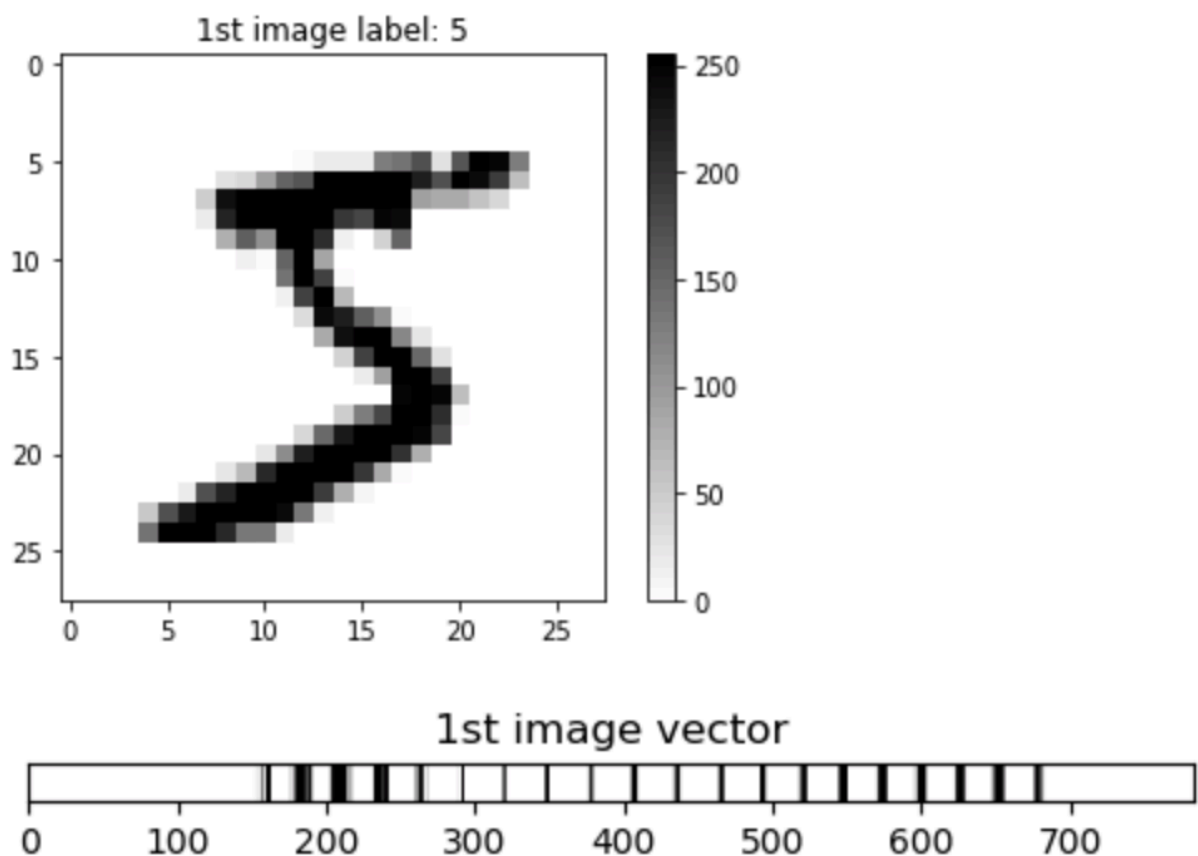


[Intro to AI] HW 3

Due on Apr. 24 at 1:00 pm

Download MNIST dataset by running *hw3_template.ipynb* from the piazza. The *train_dataset* contains a set of images and its target labels of handwritten digits. Each matrix contains a 28x28 pixel grayscale image. Here, it is assumed that you installed [PyTorch](#), but you can download MNIST from some other deep learning platforms like [Tensorflow](#).

0. **Data type conversion** This step has been implemented in *hw3_template.ipynb* and just execute the cell. Step 0 helps us work with the numpy array X (training data) and y (target labels) from now on, which is converted from torch.Tensor data type.
1. **Data visualization** We'd like to visualize the handwritten digits. Use `matplotlib.pyplot.imshow` to display the first image of X . Try displaying the same image in a vector form by using `numpy.reshape`. The images below are based on the following colormap: `imshow(..., cmap='gray_r')`




2. Principal Component Analysis (PCA)

- Vectorize all images X by using `numpy.reshape`. The shape of X is now (# of samples) by (784).

- Compute the mean μ of all the images and subtract it from X along the feature dimensions. Set X_0 to the zero-mean data matrix.
- Compute the covariance Σ of X_0 . You may use `numpy.cov` or code up the covariance.
- Compute the eigenvalues & eigenvectors of the covariance by using `np.linalg.svd`.

3. **Fraction of variance** We'd like to understand the amount of variance captured by the orthogonal directions we found in step 2.

- Make a plot of fraction of variance as a function of the number of principal components. You may use `numpy.cumsum`.



- How many principal components do we need to capture 80% of the total variance at least? Use `numpy.where` to estimate the number, and print out the number as follows:

```
We need at least ## principal components to capture 80% of the total variance.
```


4. Low-dimensional reconstruction

Reconstruct a low-dimensional version of the first image (the first row of the images matrix) from the first 100 eigenvectors. The resulting vector should be $\hat{\mathbf{x}}_i = \bar{\mathbf{x}} + \sum_{k=1}^K (\mathbf{x}_i \cdot \mathbf{e}_k) \mathbf{e}_k$, where

\mathbf{x}_i is the i -th image sample. \mathbf{e}_k is the eigenvector of k -th largest eigenvalue. $\bar{\mathbf{x}}$ is the mean of entire images. $\hat{\mathbf{x}}_i$ is the reconstructed image from K eigenvectors.

5. Visualize average image

Display the first image together with the average image μ .



6. Visualize the reconstructed image

Display the 100 images reconstructed from $\{e_1\}$, $\{e_1, e_2\}$, $\dots, \{e_1, e_2, \dots, e_{100}\}$. Use `numpy.reshape` and `imshow(..., cmap='gray_r')` as in step 1.

